

COL870: Deep Learning

Assignment 1

Harman Singh 2018EE10542

Shubham Mittal 2018EE10957

Link to Trained Models

https://drive.google.com/drive/folders/1CCIOSUgHEb7kAarZ9cKlt_7MXdCxXU_B?usp=sharing

Q1 ResNets

1.1 Image Classification using Residual Network


For this part a 14 layer ResNet model with batchnorm layers (of pytorch) is trained and validated on 40k and 10k samples from CIFAR10 train dataset respectively. Further experimental details, which follow for **Q1.2** also, are:

1. The model is trained for 100 epochs, starting with a learning rate of 0.1 which is decreased by a factor of 10 at 75th and 95th epoch.
2. Cross entropy loss and SGD optimizer is used with a batch size of 128.
3. For data augmentation on training data, random crop and random horizontal flip is used (as per the paper)

Statistics on Train, Val and Test splits :

	Accuracy	Micro F1	Macro F1
Train	0.970	0.970	0.970
Validation	0.881	0.881	0.880
Test	0.879	0.879	0.879


Accuracy Curves for train and val data



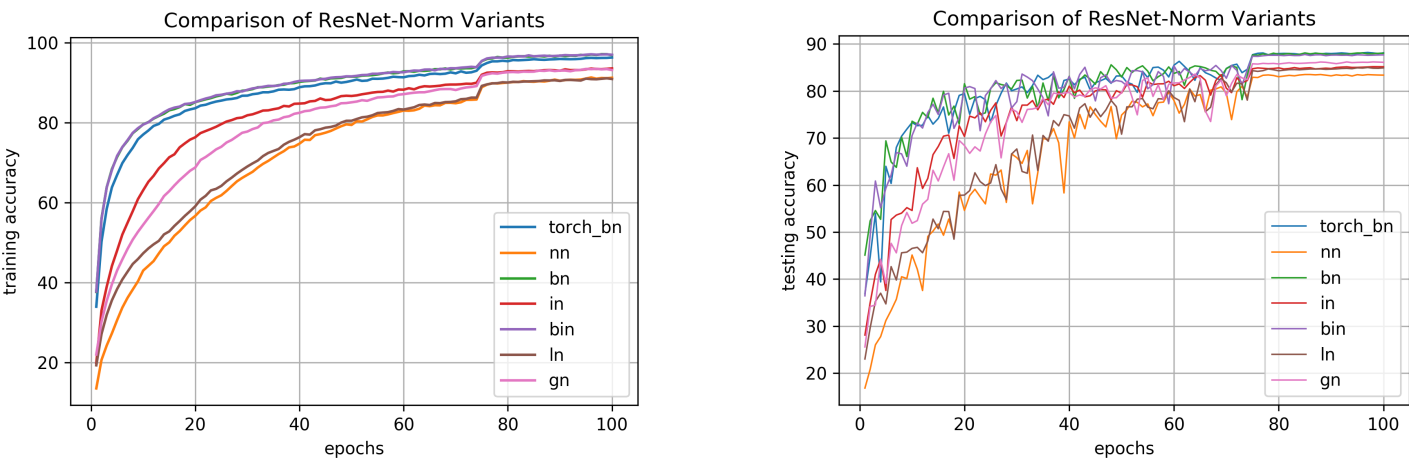
1.2 Impact of Normalization

In this part, various normalization schemes are used (implemented from scratch in pytorch). Like **Q1.1** the experimental details are the same (learning rate schedule, optimizer, loss function, batch size, etc).

1. **Comparison of torch Batch Norm and our Batch Norm** is shown below. We can see the plots are almost identical, also the validation accuracies are the same. The performance statistics of both are given along with other norm-variants.



2. **Comparing all 7 variants of ResNet Models** : torch batch norm, no norm, batch norm, instance norm, batch-instance norm, layer norm, and group norm.



Variant	Train			Validation			Test		
	Accuracy	MicroF1	MacroF1	Accuracy	MicroF1	MacroF1	Accuracy	MicroF1	MacroF1
Torch_BN	0.970	0.970	0.970	0.881	0.881	0.880	0.879	0.879	0.879
NN	0.915	0.915	0.915	0.833	0.833	0.849	0.849	0.849	0.848

BN	0.977	0.977	0.977	0.880	0.880	0.880	0.878	0.878	0.877
IN	0.923	0.923	0.923	0.851	0.851	0.851	0.847	0.847	0.847
BIN	0.977	0.977	0.977	0.876	0.876	0.877	0.881	0.881	0.881
LN	0.911	0.911	0.911	0.849	0.849	0.849	0.849	0.849	0.848
GN	0.937	0.937	0.937	0.861	0.861	0.861	0.858	0.858	0.858

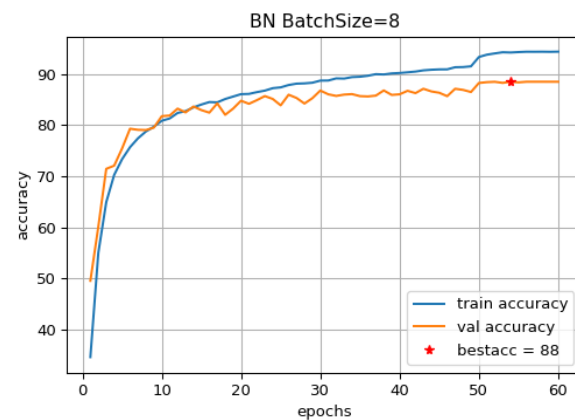

Experimental detail: Group norm variant requires another hyperparameter of defining the number of groups. We choose the number of groups in a norm layer, accepting n number of channels, as $n/4$. For example - for norm layer taking input of 16 channels, it uses 4 groups for normalisation.


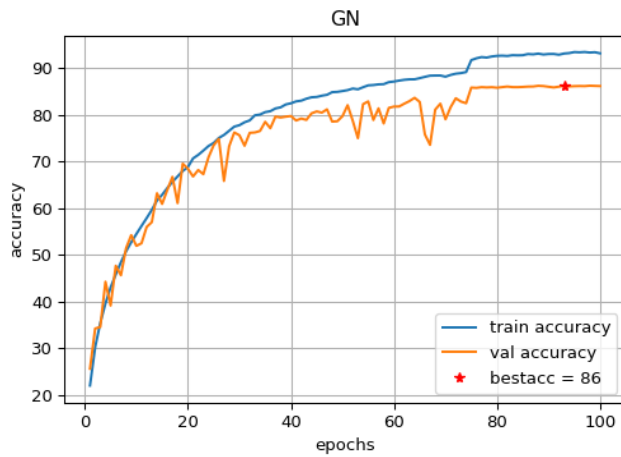
Inferences:

1. The accuracy plots of torch_BN and BN variants of ResNet are overlapping. Also the performance statistics are almost same. Hence the sanity check (1.2.4) is done and verified.
2. From the plots, we can see that BIN learns most rapidly i.e. in about 10 epochs it has attained about 80% train acc and 70% val acc, whereas in same number of epochs the NN variant is at 40% train and val acc.
3. The performance statistics also conclude that the BIN variant performs the best.

3. Impact of Batch Size

Batch normalisation and group normalisation are compared: left- batch size = 128, right- batch size=8





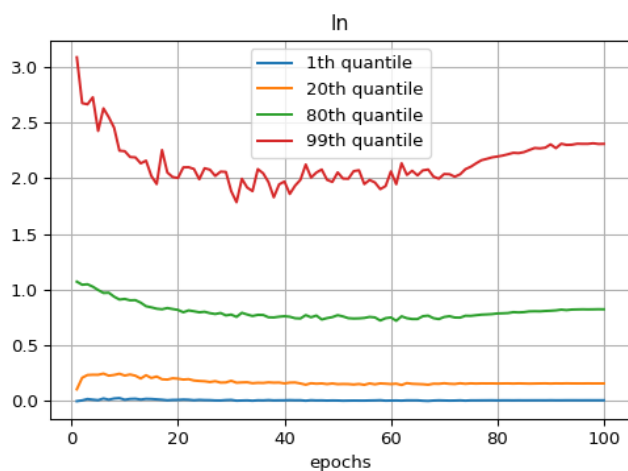






Inferences/Comments:

1. The training time substantially increased by almost 6 times when batch size reduced from 128 to 8.
2. To save time, we early stop at the 60th epoch in right plots.
3. The learning curves clearly show that Group norm is almost invariant to batch size. To our surprise, the batchnorm is also showing no degradation in performance even when batch size is reduced. This needs further investigation but our intuitions says that when the batch size is less, the learning will take more time, but the dataset is probably such that the convergence has taken place even with small batch size, and this convergence may have taken place may have been at a similar position in the loss space giving similar accuracy. This needs further investigation

4. Evolution of feature distributions





Inferences

1. From the above plots, we can infer/see how the models can be trained “healthily” using norm layers.
2. NN variant features are diverging with training whereas torch_bn, BN, BIN variants are stable with time.

- Thus, we conclude that performing normalization is also helpful in controlling the distribution of features (as discussed in the group normalization paper).

Q2 LSTM with Layer Normalization and CRF

2.1 NER Tagging with Bi-LSTM

2.1.1 Experimental Details

The following hyperparameters were found to be optimal as per our search. Unless specified otherwise, we use these to train all our models

- Adam optimizer with initial learning rate of 0.001, trained for 50 epochs. We do not see overfitting when the model is trained for a long time (>50 epochs).
- Dropout prob = 0.5, BatchSize = 128, gradient clipping at 5 as given in the assignment.
- Scheduling - we schedule the learning such that at 35 epochs, we change the optimizer to Adam with initial learning rate of 0.0001 just so that model can converge better, but did not help improve the accuracy much in our case.
- Loss function used is cross entropy except in the case of BiLSTM-CRF. Experiments with weighted cross entropy are not reported since there was no major gain in performance, in fact the performance was similar with and without weighting, if trained for long enough (50 epochs in our case).

For the plots of performance statistics, for each epoch, we calculate the average of accuracy, micro f1 and macro f1 score, **at a token level excluding “O” tags and “pad” tags**, across all batches and save it to plot after training


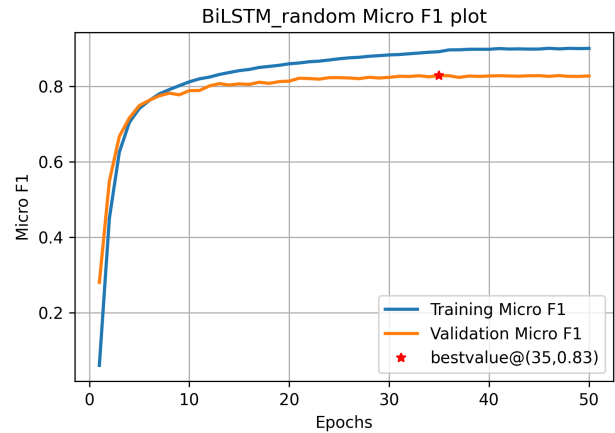

2.1.2 Simple BiLSTM results

- **BiLSTM with Random Word Embeddings**

CONLL based final performance metrics (after training):

	Accuracy	Micro F1	Macro F1
Train	0.982	0.889	0.755
Validation	0.964	0.803	0.619
Test	0.964	0.804	0.614

Plots for performance statistics vs epochs:

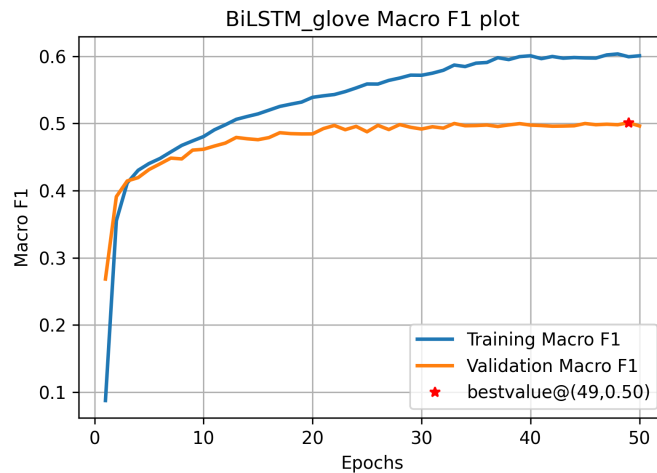

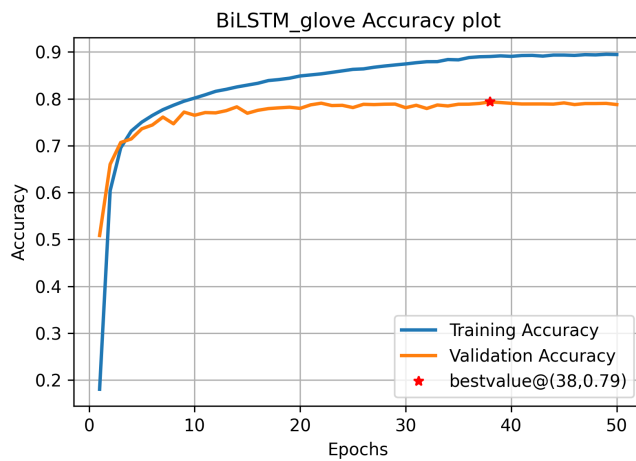


- BiLSTM with Pretrained (Glove) word embeddings

CONLL based final performance metrics (after training):


	Accuracy	Micro F1	Macro F1
Train	0.986	0.909	0.823
Validation	0.961	0.793	0.620
Test	0.962	0.794	0.614

Plots for performance statistics vs epochs:



- Comparison of Simple BiLSTM with Random and Glove embeddings

(Following Plots, 3 rows respectively are for accuracy, microF1 and macroF1. Left column - Training, Right column - Validation)
(orange line - glove, blue line - random)



Comments

- We see that with glove embeddings the model trains faster than with random embeddings (ie with glove embeddings it takes less number of epochs for same performance). We conclude this because for all the plots, accuracy, microF1 and MacroF1, for both train and validation sets, the curve for BiLSTM with glove embeddings remains above that of random embeddings
- Faster training is expected because glove embeddings provide embeddings, which are similar for similar words, and this pretraining of the embeddings means that we don't have to learn the embedding now, we can just fine tune it during training. While in case of random embeddings all the embeddings also have to be learnt which takes more time.

- The final performance of both BiLSTM's are similar hence we can say that random embeddings are changed during training (learnt) to represent meaningful information from words just like glove embeddings.

2.1.3 Results using BiLSTM using character level features and pretrained embeddings



We use a separate BiLSTM for generating forward and backward character level embeddings of the Word which is concatenated with the word embeddings. Total embedding size now becomes $100+25+25 = 150$


- BiLSTM with Character Level Embeddings + Pretrained (Glove) word embeddings (concatenated)

CONLL based final performance metrics (after training):

	Accuracy	Micro F1	Macro F1
Train	0.986	0.908	0.820
Validation	0.961	0.795	0.629
Test	0.962	0.795	0.626

Plots for performance statistics vs epochs:





- **Comparison between BiLSTM (Char Level + Glove) model with Bilstm (Glove)**







On the basis of Final Performance:

Model	Validation		Test	
	MicroF1	MacroF1	MicroF1	MacroF1
BiLSTM (Glove embeds)	0.793	0.620	0.794	0.614
BiLSTM (Char Level + Glove embeds)	0.795	0.629	0.795	0.626

On the basis of learning curves:

(Following Plots, 3 rows respectively are for accuracy, microF1 and macroF1. Left column - Training, Right column - Validation)

(orange line - glove, blue line - char+glove)



Comments :

- For Validation and Test data, the BiLSTM with character and word embeddings perform better in terms of both micro and macro F1 score. This is expected since we are giving the model character level information of the words, hence this will help enhancing the embeddings of the

words (giving more information) and would also be helpful to words which are not in the training set (out of vocabulary), to have better representation than the standard “<OOV>” tag during validation and test time.

- The training curves for both models are very similar, hence speed of learning for both the models is almost the same, ie the additional lstm required for getting the character level embeddings does not consume more time (this is due to fully vectorized implementation to get char level embeddings without looping over words).

2.1.4 Results using BiLSTM with LayerNormalization, using character level features and pretrained embeddings

- Implementation details of LSTM with Layer Normalization:

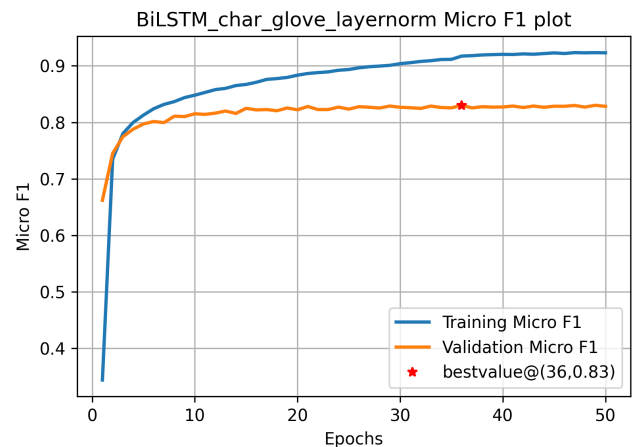
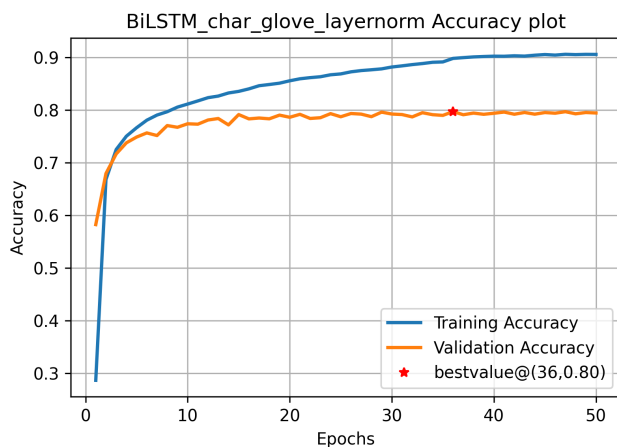
1. We replace nn.LSTM with our implementation of LSTM, which uses layer-normed LSTM Cell implementation.
2. Inside the LSTM cell, since there are 5 non-linearities, thus the input to them is normalized.
3. In total, three layer norm layers are used. Two are applied on the linear transformations : $W_{x2h} \cdot x_t$ and $W_{h2h} \cdot h_{t-1}$. Third layer norm layer is applied on c_t .


- BiLSTM with Layer Normalization (used Character Level Embeddings + Pretrained (Glove) word embeddings (concatenated))

CONLL based final performance metrics (after training):

	Accuracy	Micro F1	Macro F1
Train	0.989	0.923	0.851
Validation	0.963	0.799	0.625
Test	0.963	0.802	0.627

Plots for performance statistics vs epochs:





- Comparison between BiLSTM (Char Level + Glove) with and without Layer Normalization




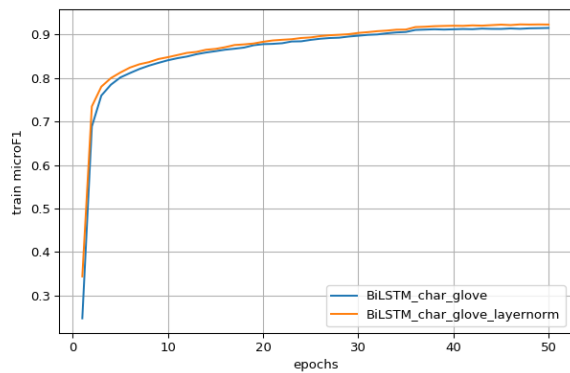
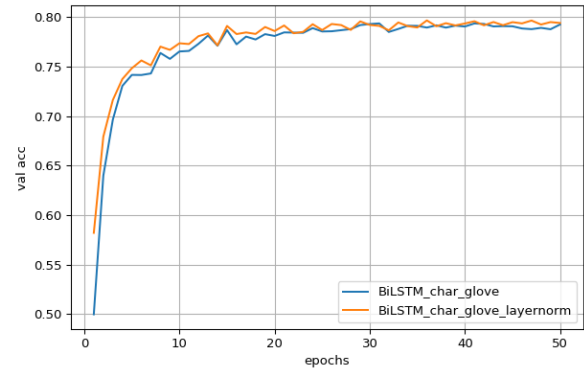

On the basis of Final Performance:

Model	Validation		Test	
	MicroF1	MacroF1	MicroF1	MacroF1
BiLSTM (Glove embeds)	0.793	0.620	0.794	0.614
BiLSTM (Char Level + Glove embeds)	0.795	0.629	0.795	0.626
Without Layer Norm				
BiLSTM (Char Level + Glove embeds)	0.799	0.625	0.802	0.627
With Layer Norm				

On the basis of learning curves:

(Following Plots, 3 rows respectively are for accuracy, microF1 and macroF1. Left column - Training, Right column - Validation)

(orange line - char+glove+layernorm, blue line - char+glove)



Comments :

- For Validation and Test data, the BiLSTM with Layer Normalization performs better in terms of Micro F1 score and similar in terms of macro F1 score on test data.
- In most of the curves above, normalized variant of the model, has its curve above the unnormalized. Hence speed of training with layer normalization is more that without it. This is expected because loss landscape with normalization becomes smoother and even (in terms of

gradient) in all directions which lets the model take larger gradient steps (in case of adam optimizer which we have used) hence faster training.

2.2 Linear chain CRF

2.2.2

Initialization:

- The appropriate initialization of the transition matrix would be to choose small random values for all the cells of the transition matrix **except the cells representing transition from stop state, or transition to start state**. These cells should be initialized using a very high negative value (like -10000, -20000) so that the exponent of that tends to 0 basically suggesting that these transitions are impossible(logically we cannot transition to start state or transition from stop state, these have to be beginning and end states respectively).
- Another intuition is that if the forbidden transitions take place then, the large negative value of the transition score would make the exponent in the numerator of the likelihood function nearly 0, which would cause the negative log likelihood to go to +inf. To avoid this, these transitions would never take place.

Loss function for CRF:

- The negative of the log likelihood function described in the paper was used as our loss function.

2.2.3 BiLSTM + CRF layer results

We experimented by first using out BiLSTM + CRF model with character level features and pretrained word embeddings which gave us a performance boost. After this we also tried out adding layer norm to the above architecture (using layer normalized BiLSTM) which further improved our performance.




Results for both experiments are shown below

- **BiLSTM (Char level features + pretrained embeddings) + CRF**

CONLL based final performance metrics (after training):

	Accuracy	Micro F1	Macro F1
Train	0.988	0.930	0.890
Validation	0.962	0.814	0.643
Test	0.962	0.817	0.634




Plots for performance statistics vs epochs:



- **BiLSTM with Layer Normalization (Char level features + pretrained embeddings)+ CRF**
CONLL based final performance metrics (after training):

	Accuracy	Micro F1	Macro F1
Train	0.990	0.941	0.907
Validation	0.964	0.820	0.649
Test	0.964	0.820	0.635

Plots for performance statistics vs epochs:



- Comparison between BiLSTM with CRF with other variants

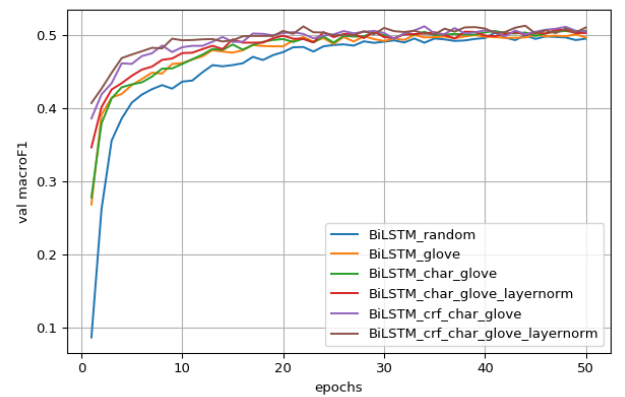





On the basis of Final Performance:

Model	Validation		Test	
	MicroF1	MacroF1	MicroF1	MacroF1
BiLSTM (random embeddings)	0.803	0.619	0.804	0.614
BiLSTM (Glove embeddings)	0.793	0.620	0.794	0.614
BiLSTM (Glove + char embeddings)	0.795	0.629	0.795	0.626

BiLSTM with LayerNorm (Glove + char embeddings)	0.799	0.625	0.802	0.627
BiLSTM with CRF (Glove + char embeddings)	0.814	0.643	0.817	0.634
BiLSTM with LayerNorm + CRF (Glove + char embeddings)	0.820	0.649	0.820	0.635

On the basis of learning curves:

(Following Plots, 3 rows respectively are for accuracy, microF1 and macroF1. Left column - Training, Right column - Validation)



Comments :

- The BiLSTM with CRF layer implemented on top of it, gives the best performance among all other variants for NER tagging. The gain in performance is also the highest, as compared to any other performance enhancer (layernorm/char embeddings/glove embeddings).

- This can be attributed to CRF being able to capture relationship between words using the transition scores internally, which can be thought of as a regularizer, which imposes (and learns) certain priors on the output sentences (whether one word can follow the other, etc)
- We see a gradual increase in performance as we increase the complexity of the model and introduce new performance enhancing measures. Glove embeddings proved to be better than random embeddings in terms of learning curves (faster learning). Including character level features improved performance and gave the model better representation of the tokens. This was improved upon by the normalization by the layernorm. Finally the CRF model with layer norm proved to be the best for ner tagging in our case.
- Referring to the curves above, we see that Bilstm with layernorm, crf and using char and glove embeddings learns the fastest among all models in general. The layernorm has its advantage in speed of training as described in the layernorm section of this document, and crf introduces extra constraints and learns the relationship between words in a sentence which may have helped us achieve faster training.
- Without layer norm, Bilstm with crf and using char and glove has slightly less speed of training and has the curve below that of the layernorm variant.
- The general trend of the plots is that the speed of training increases with increase in complexity of the model, and the complexity can be introduced by adding layernorm, crf, glove embeddings, char level embeddings etc. Best model is the combination of all, both in terms of speed and final accuracy.

Appendix

2.1.2

- Classification Report for Simple BiLSTM with Random Embeddings
(Train, Validation and Test data respectively)

PERFORMANCE ON Train DATA

MicroF1 = 0.8891281071185972

Accuracy = 0.9821888393554009

-----Classification Report-----

	precision	recall	f1-score	support
art	0.673	0.341	0.453	296
eve	0.543	0.478	0.508	226
geo	0.886	0.937	0.911	29240
gpe	0.968	0.942	0.955	12058
nat	0.714	0.564	0.630	133
org	0.805	0.777	0.791	15803
per	0.879	0.879	0.879	13121
tim	0.922	0.911	0.917	15767
micro avg	0.887	0.891	0.889	86644
macro avg	0.799	0.729	0.755	86644
weighted avg	0.886	0.891	0.888	86644

PERFORMANCE ON Validation DATA

MicroF1 = 0.8034243979743

Accuracy = 0.9635724503487508

-----Classification Report-----

	precision	recall	f1-score	support
art	0.265	0.086	0.129	105
eve	0.380	0.244	0.297	78
geo	0.823	0.866	0.844	9724
gpe	0.936	0.908	0.922	4210
nat	0.629	0.440	0.518	50
org	0.635	0.623	0.629	5187
per	0.757	0.744	0.750	4457
tim	0.881	0.853	0.867	5254
micro avg	0.804	0.802	0.803	29065
macro avg	0.663	0.595	0.619	29065
weighted avg	0.803	0.802	0.802	29065

PERFORMANCE ON Test DATA

MicroF1 = 0.8038678760910491

Accuracy = 0.963636161631873

-----Classification Report-----

	precision	recall	f1-score	support
art	0.280	0.069	0.110	102
eve	0.378	0.322	0.348	87
geo	0.831	0.871	0.851	9912
gpe	0.928	0.907	0.918	4168
nat	0.571	0.364	0.444	55
org	0.634	0.623	0.628	5205
per	0.758	0.753	0.756	4406
tim	0.868	0.844	0.856	5275
micro avg	0.804	0.804	0.804	29210
macro avg	0.656	0.594	0.614	29210
weighted avg	0.802	0.804	0.802	29210

- Classification Report for Simple BiLSTM with Glove

(Train, Validation and Test data respectively)

PERFORMANCE ON Train DATA

MicroF1 = 0.9092593551654092

Accuracy = 0.9859880455904748

-----Classification Report-----

	precision	recall	f1-score	support
art	0.782	0.632	0.699	296
eve	0.716	0.681	0.698	226
geo	0.901	0.949	0.925	29240
gpe	0.970	0.952	0.961	12058
nat	0.728	0.564	0.636	133
org	0.848	0.809	0.828	15803
per	0.908	0.902	0.905	13121
tim	0.930	0.935	0.933	15767
micro avg	0.907	0.912	0.909	86644
macro avg	0.848	0.803	0.823	86644
weighted avg	0.906	0.912	0.909	86644

PERFORMANCE ON Validation DATA

MicroF1 = 0.7928936234766631

Accuracy = 0.9614248697959983

-----Classification Report-----

	precision	recall	f1-score	support
art	0.500	0.152	0.234	105
eve	0.400	0.282	0.331	78
geo	0.807	0.872	0.838	9724
gpe	0.934	0.914	0.924	4210
nat	0.514	0.380	0.437	50
org	0.583	0.638	0.609	5187
per	0.737	0.735	0.736	4457
tim	0.844	0.866	0.855	5254
micro avg	0.777	0.809	0.793	29065
macro avg	0.665	0.605	0.620	29065
weighted avg	0.779	0.809	0.793	29065

PERFORMANCE ON Test DATA

MicroF1 = 0.7943733843623058

Accuracy = 0.9621140577955047

-----Classification Report-----

	precision	recall	f1-score	support
art	0.480	0.118	0.189	102
eve	0.341	0.322	0.331	87
geo	0.818	0.875	0.846	9912
gpe	0.931	0.916	0.923	4168
nat	0.621	0.327	0.429	55
org	0.585	0.634	0.609	5205
per	0.735	0.743	0.739	4406
tim	0.836	0.860	0.848	5275
micro avg	0.779	0.810	0.794	29210
macro avg	0.668	0.599	0.614	29210
weighted avg	0.780	0.810	0.794	29210

2.1.3

- Classification Report for Simple BiLSTM with Character + Word embeddings

(Train, Validation and Test data respectively)

PERFORMANCE ON Train DATA

MicroF1 = 0.9083878233489238

Accuracy = 0.9860040241349769

-----Classification Report-----

	precision	recall	f1-score	support
art	0.738	0.581	0.650	296
eve	0.738	0.673	0.704	226
geo	0.899	0.952	0.925	29240
gpe	0.972	0.949	0.960	12058
nat	0.741	0.602	0.664	133
org	0.850	0.816	0.832	15803
per	0.905	0.896	0.900	13121
tim	0.923	0.934	0.928	15767
micro avg	0.904	0.912	0.908	86644
macro avg	0.846	0.800	0.820	86644
weighted avg	0.904	0.912	0.908	86644

PERFORMANCE ON Validation DATA

MicroF1 = 0.7951207958313595

Accuracy = 0.9612918011214732

-----Classification Report-----

	precision	recall	f1-score	support
art	0.370	0.162	0.225	105
eve	0.469	0.295	0.362	78
geo	0.809	0.875	0.841	9724
gpe	0.938	0.910	0.924	4210
nat	0.512	0.420	0.462	50
org	0.563	0.640	0.599	5187
per	0.774	0.729	0.751	4457
tim	0.870	0.862	0.866	5254
micro avg	0.782	0.808	0.795	29065
macro avg	0.663	0.612	0.629	29065
weighted avg	0.786	0.808	0.796	29065

```

PERFORMANCE ON Test DATA
MicroF1 = 0.7945924132364811
Accuracy = 0.9616400205909911
-----Classification Report-----

```

	precision	recall	f1-score	support
art	0.308	0.118	0.170	102
eve	0.444	0.368	0.403	87
geo	0.818	0.880	0.848	9912
gpe	0.931	0.911	0.921	4168
nat	0.600	0.382	0.467	55
org	0.564	0.636	0.598	5205
per	0.766	0.734	0.749	4406
tim	0.852	0.853	0.852	5275
micro avg	0.781	0.809	0.795	29210
macro avg	0.660	0.610	0.626	29210
weighted avg	0.784	0.809	0.795	29210

2.1.4

- Classification Report for BiLSTM with Layer Normalization using Character + Word embeddings (Train, Validation and Test data respectively)

PERFORMANCE ON Train DATA

MicroF1 = 0.9233053568152632

Accuracy = 0.9887547920269522

-----Classification Report-----

	precision	recall	f1-score	support
art	0.779	0.713	0.744	296
eve	0.683	0.752	0.716	226
geo	0.915	0.958	0.936	29240
gpe	0.972	0.955	0.964	12058
nat	0.673	0.789	0.727	133
org	0.875	0.848	0.861	15803
per	0.918	0.917	0.917	13121
tim	0.931	0.955	0.943	15767
micro avg	0.917	0.929	0.923	86644
macro avg	0.843	0.861	0.851	86644
weighted avg	0.917	0.929	0.923	86644

PERFORMANCE ON Validation DATA

MicroF1 = 0.7989299742652038

Accuracy = 0.9626003097543034

-----Classification Report-----

	precision	recall	f1-score	support
art	0.385	0.190	0.255	105
eve	0.289	0.282	0.286	78
geo	0.810	0.874	0.841	9724
gpe	0.940	0.912	0.926	4210
nat	0.460	0.460	0.460	50
org	0.597	0.643	0.620	5187
per	0.767	0.736	0.752	4457
tim	0.850	0.870	0.860	5254
micro avg	0.786	0.812	0.799	29065
macro avg	0.637	0.621	0.625	29065
weighted avg	0.788	0.812	0.799	29065

PERFORMANCE ON Test DATA

MicroF1 = 0.801542399137873

Accuracy = 0.9631102766081157

-----Classification Report-----

	precision	recall	f1-score	support
art	0.245	0.118	0.159	102
eve	0.361	0.402	0.380	87
geo	0.822	0.881	0.850	9912
gpe	0.936	0.910	0.923	4168
nat	0.535	0.418	0.469	55
org	0.598	0.645	0.621	5205
per	0.764	0.746	0.755	4406
tim	0.845	0.866	0.855	5275
micro avg	0.789	0.815	0.802	29210
macro avg	0.638	0.623	0.627	29210
weighted avg	0.790	0.815	0.802	29210

2.2

- Classification Report for BiLSTM + CRF with Character + Word embeddings (Train, Validation and Test data respectively)

PERFORMANCE ON Train DATA

MicroF1 = 0.93000374866634

Accuracy = 0.9880861513954801

-----Classification Report-----

	precision	recall	f1-score	support
art	0.912	0.767	0.833	296
eve	0.846	0.827	0.837	226
geo	0.923	0.958	0.940	29240
gpe	0.975	0.961	0.968	12058
nat	0.775	0.827	0.800	133
org	0.887	0.850	0.868	15803
per	0.934	0.920	0.927	13121
tim	0.947	0.951	0.949	15767
micro avg	0.929	0.931	0.930	86644
macro avg	0.900	0.883	0.890	86644
weighted avg	0.929	0.931	0.930	86644

PERFORMANCE ON Validation DATA

MicroF1 = 0.8143543471148381

Accuracy = 0.9615875092870846

-----Classification Report-----

	precision	recall	f1-score	support
art	0.500	0.171	0.255	105
eve	0.423	0.282	0.338	78
geo	0.828	0.870	0.848	9724
gpe	0.939	0.921	0.930	4210
nat	0.524	0.440	0.478	50
org	0.645	0.656	0.650	5187
per	0.803	0.740	0.770	4457
tim	0.884	0.866	0.875	5254
micro avg	0.815	0.814	0.814	29065
macro avg	0.693	0.618	0.643	29065
weighted avg	0.815	0.814	0.813	29065

PERFORMANCE ON Test DATA

MicroF1 = 0.816718054128126

Accuracy = 0.9620807270545624

-----Classification Report-----

	precision	recall	f1-score	support
art	0.414	0.118	0.183	102
eve	0.417	0.345	0.377	87
geo	0.838	0.877	0.857	9912
gpe	0.931	0.917	0.924	4168
nat	0.478	0.400	0.436	55
org	0.651	0.656	0.654	5205
per	0.801	0.748	0.774	4406
tim	0.879	0.862	0.870	5275
micro avg	0.817	0.816	0.817	29210
macro avg	0.676	0.615	0.634	29210
weighted avg	0.816	0.816	0.816	29210

- Classification Report for BiLSTM + CRF with Character + Word embeddings AND Layer Norm (Train, Validation and Test data respectively)

PERFORMANCE ON Train DATA

MicroF1 = 0.9408832687736344

Accuracy = 0.9901621330619439

-----Classification Report-----

	precision	recall	f1-score	support
art	0.900	0.818	0.857	296
eve	0.858	0.832	0.845	226
geo	0.933	0.968	0.950	29240
gpe	0.977	0.967	0.972	12058
nat	0.800	0.902	0.848	133
org	0.919	0.869	0.893	15803
per	0.934	0.937	0.936	13121
tim	0.944	0.964	0.954	15767
micro avg	0.938	0.943	0.941	86644
macro avg	0.908	0.907	0.907	86644
weighted avg	0.938	0.943	0.941	86644

PERFORMANCE ON Validation DATA

MicroF1 = 0.8200622621930128

Accuracy = 0.964174955736184

-----Classification Report-----

	precision	recall	f1-score	support
art	0.346	0.171	0.229	105
eve	0.450	0.346	0.391	78
geo	0.834	0.868	0.851	9724
gpe	0.931	0.923	0.927	4210
nat	0.480	0.480	0.480	50
org	0.714	0.648	0.679	5187
per	0.784	0.754	0.769	4457
tim	0.865	0.873	0.869	5254
micro avg	0.824	0.816	0.820	29065
macro avg	0.676	0.633	0.649	29065
weighted avg	0.821	0.816	0.818	29065

PERFORMANCE ON Test DATA

MicroF1 = 0.8200697390796503

Accuracy = 0.9640398339388417

-----Classification Report-----

	precision	recall	f1-score	support
art	0.233	0.098	0.138	102
eve	0.397	0.356	0.376	87
geo	0.840	0.871	0.855	9912
gpe	0.924	0.924	0.924	4168
nat	0.464	0.473	0.468	55
org	0.705	0.644	0.673	5205
per	0.784	0.767	0.775	4406
tim	0.863	0.871	0.867	5275
micro avg	0.823	0.817	0.820	29210
macro avg	0.651	0.625	0.635	29210
weighted avg	0.819	0.817	0.818	29210